

# What do you learn from context?

Probing for sentence structure in contextualized word representations.

Review by Gong Qiong

# Outline

- Abstract
- Introduction
- Edge Probing
- Experiment Design
- Results and Analysis
- Discussion

# Abstract

- Contextualized representation models such as ELMo and BERT have recently achieved state-of-the-art results on a diverse array of downstream NLP tasks.
- Building on recent token-level probing work, we introduce a novel edge probing task design and construct a broad suite of sub-sentence tasks derived from the traditional structured NLP pipeline.

# Abstract

- We probe word-level contextual representations from four recent models and investigate how they encode sentence structure across a range of syntactic, semantic, local, and long-range phenomena.
- We find that existing models trained on language modeling and translation produce strong representations for syntactic phenomena, but only offer comparably small improvements on semantic tasks over a non-contextual baseline.

# Introduction

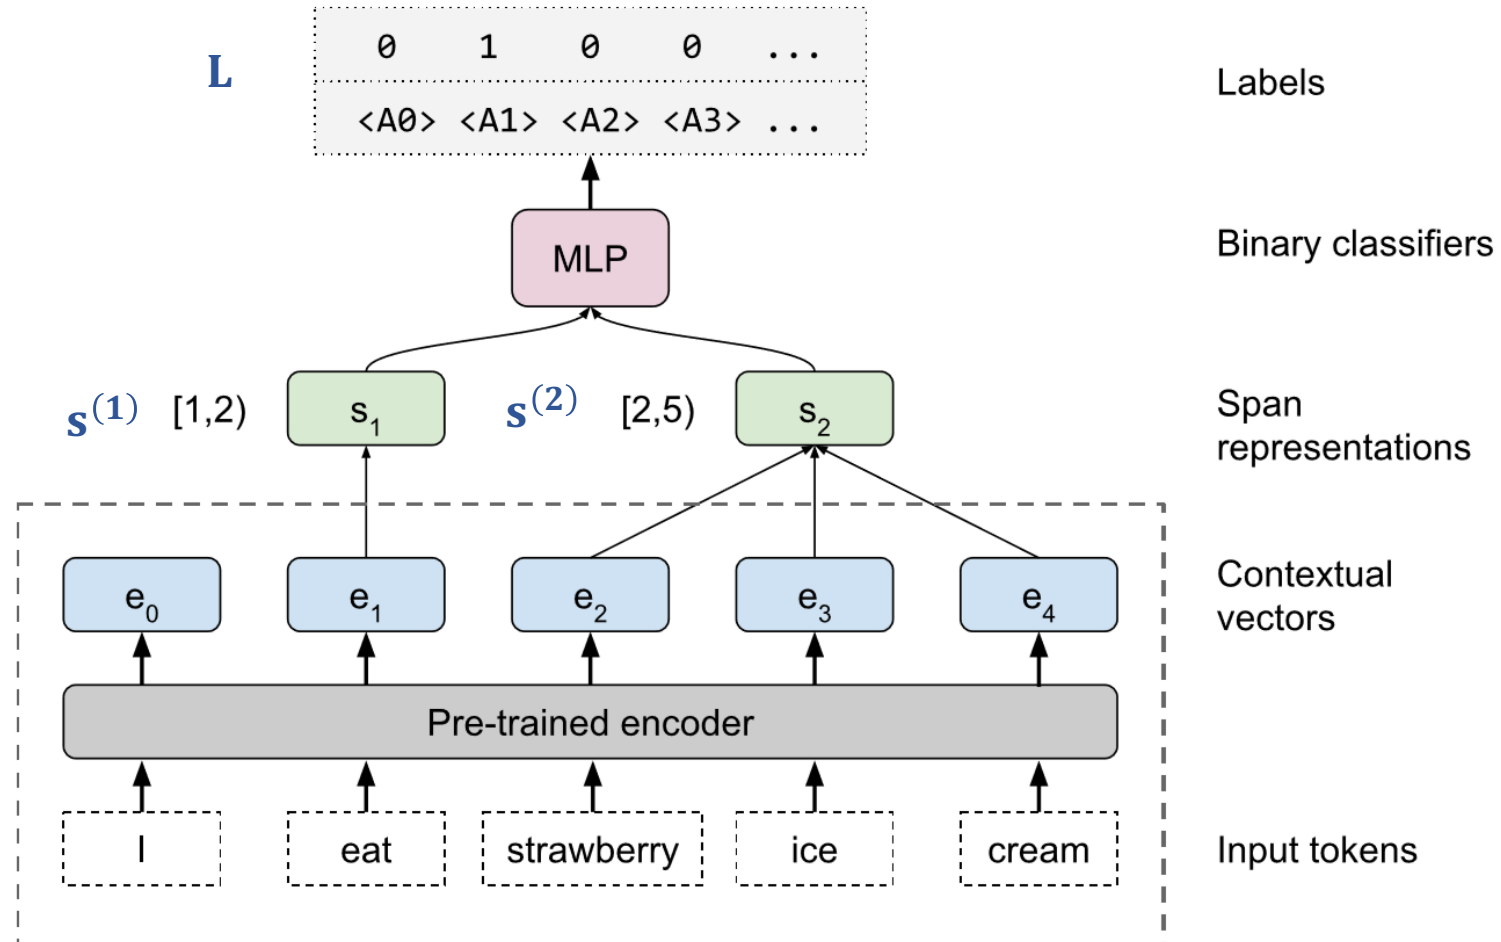
# Introduction - 2 Qs

What do you learn from context

- **Syntactic** or **Semantic**?
- **Local** or **Long-range**?

What do you learn from context?

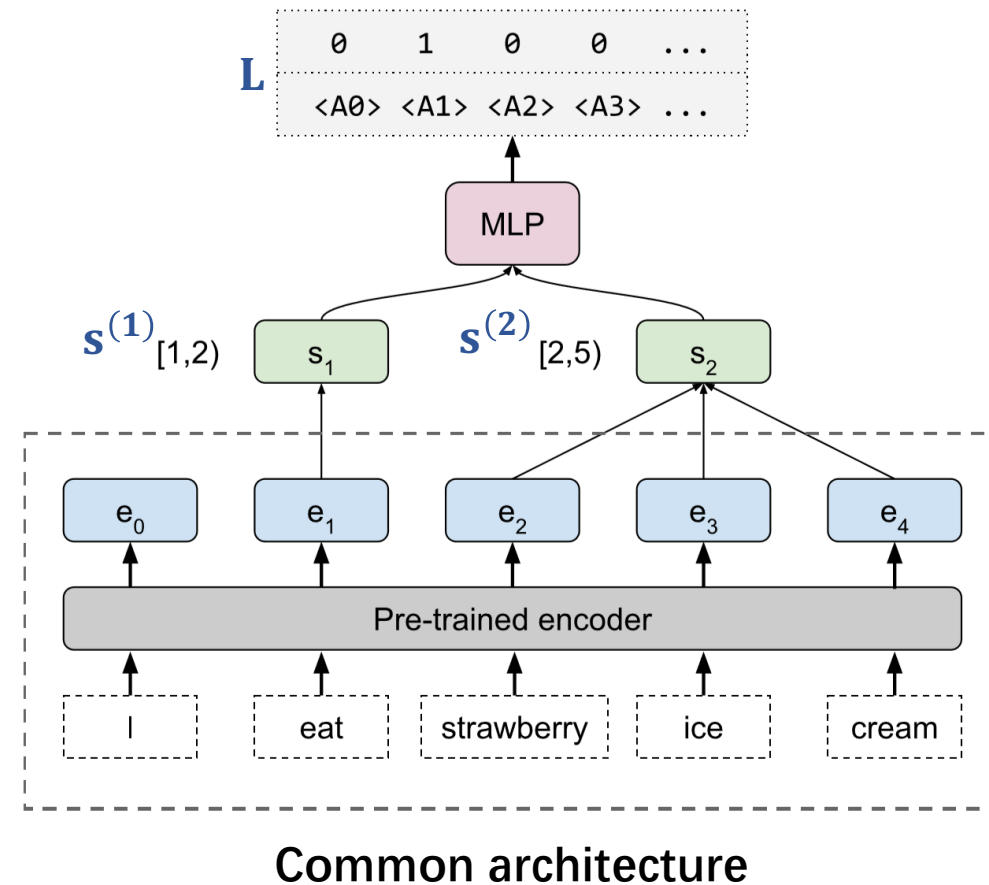
# Edge probing - model architecture



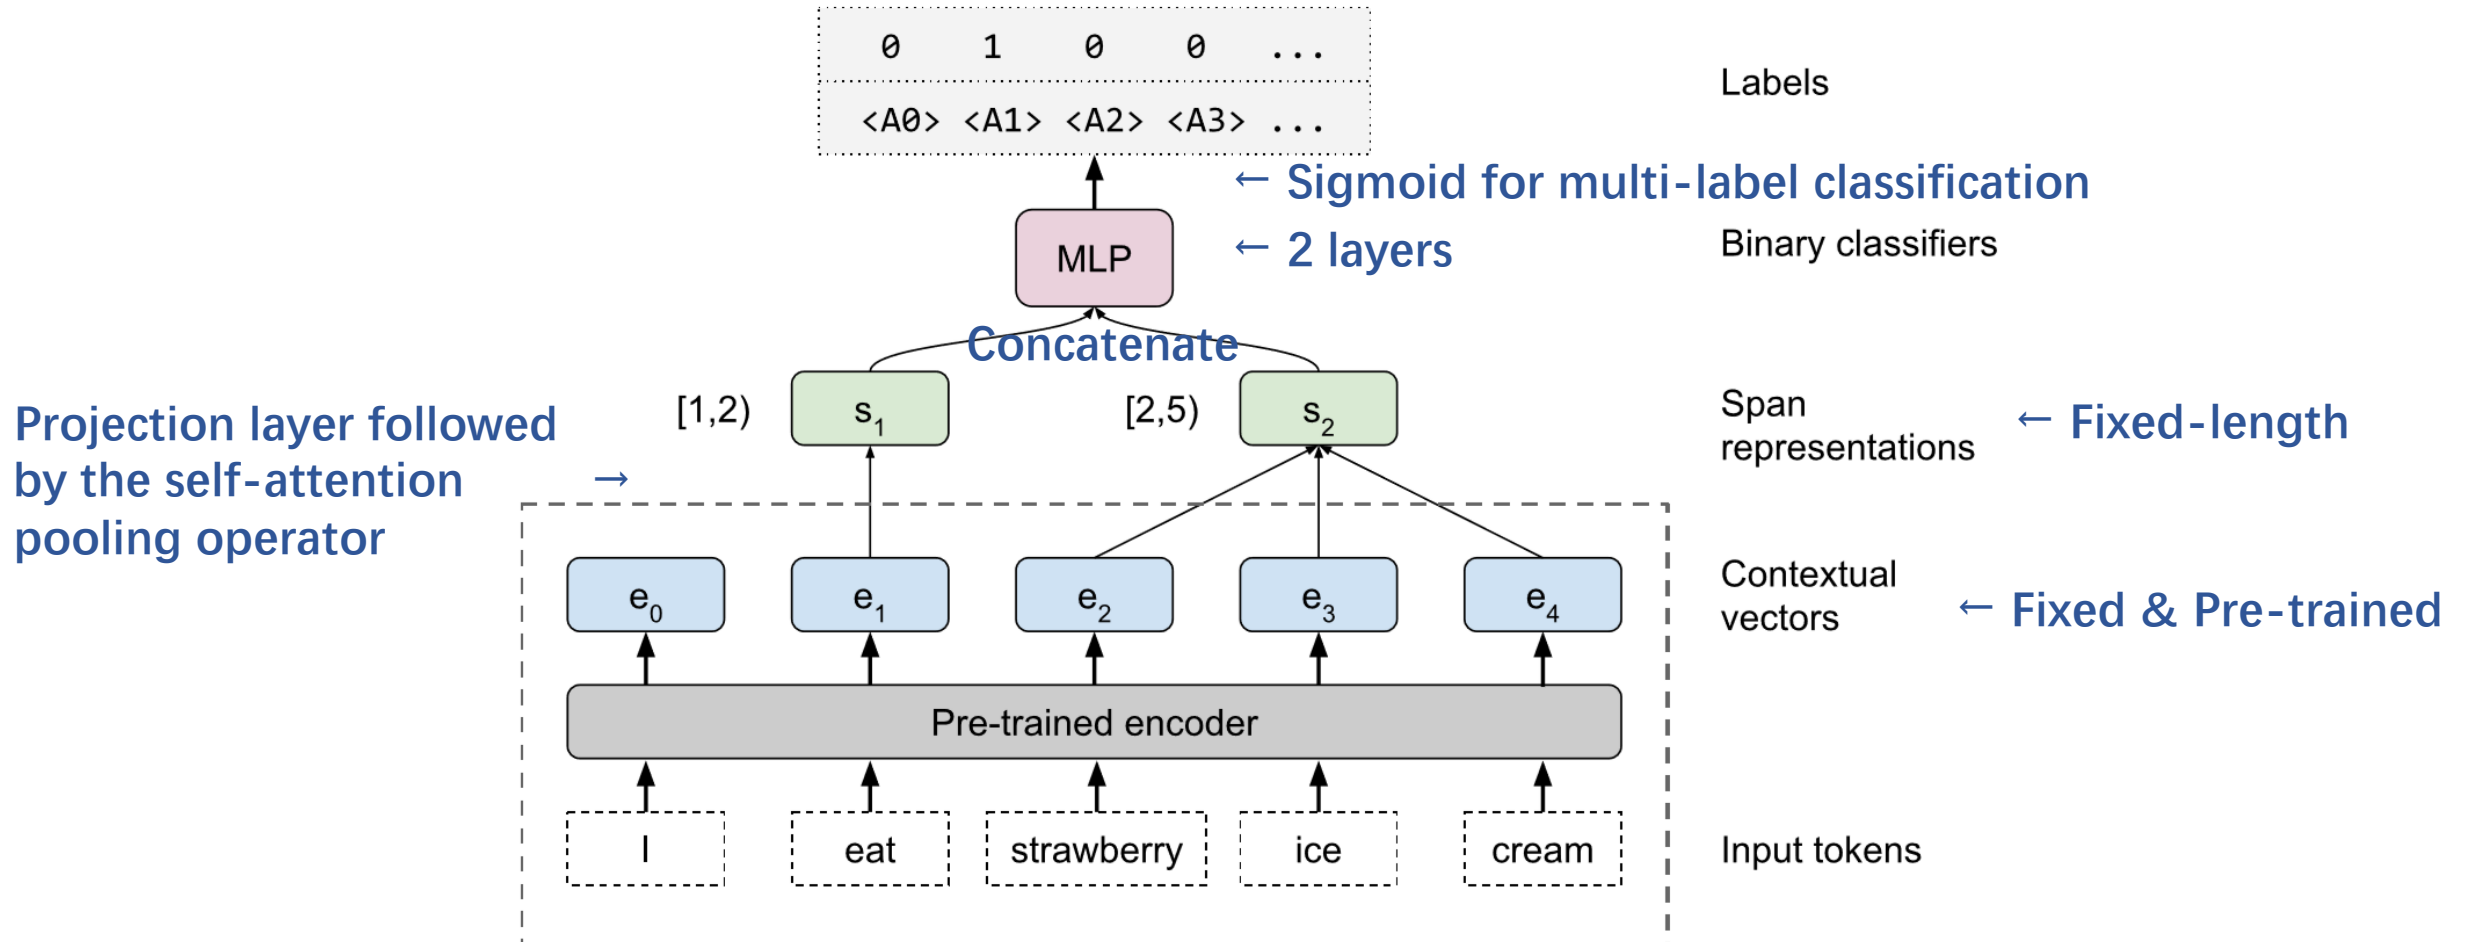


# Edge probing

- Represent a sentence as a list of tokens
  - $T = [t_0, t_1, \dots, t_n]$
- Input a labeled edge as  $\{s^{(1)}, s^{(2)}, L\}$ 
  - $s^{(1)} = [i^{(1)}, j^{(1)})$
  - $s^{(2)} = [i^{(2)}, j^{(2)})$ , optional (as Coreference), omitted for unary edges (as POS tagging)
  - $L$ : a set of one or more targets from a task-specific label set.
- Predict  $L$  as multi-label classification
  - **Uniform metric**: binary F1 score



# Edge probing - model architecture



# Edge probing

- **Separate projections**

- Why? -> Different regulations can be extracted from different parts ( $s^{(1)}, s^{(2)}$ )

- **Self-Attention Pooling**

- 1. Calculate a weight.
  - 2. Compute a weighted span representations.
  - Why? -> Strengthen the model so that it can perform better, making the result more obvious and prominent.

- **MLP: 2 layers**


# Experiment Design

# NLP Tasks

- Syntactic tasks
  - Part-of-speech tagging (POS)
  - Constituent labeling
  - Dependency labeling
  - Named entity labeling
  - Semantic role labeling (SRL) core roles
- Semantic tasks
  - Semantic proto-role (SPR)
  - Semantic role labeling (SRL) non-core roles
  - Coreference (OntoNote + Winograd-style)
  - Relation Classification (Rel.)

Various!

# NLP Tasks – example



	0	1	2	3	4	5	6	7	8	9	10	11
POS	The important thing about Disney is that it is a global [brand] <sub>1</sub> . → NN (Noun)											
Constit.	The important thing about Disney is that it [is a global brand] <sub>1</sub> . → VP (Verb Phrase)											
Depend.	[Atmosphere] <sub>1</sub> is always [fun] <sub>2</sub> → nsubj (nominal subject)											
Entities	The important thing about [Disney] <sub>1</sub> is that it is a global brand. → Organization											
SRL	[The important thing about Disney] <sub>2</sub> [is] <sub>1</sub> that it is a global brand. → Arg1 (Agent)											
SPR	[It] <sub>1</sub> [endorsed] <sub>2</sub> the White House strategy. . . → {awareness, existed_after, . . . }											
Coref. <sup>O</sup>	The important thing about [Disney] <sub>1</sub> is that [it] <sub>2</sub> is a global brand. → True											
Coref. <sup>W</sup>	[Characters] <sub>2</sub> entertain audiences because [they] <sub>1</sub> want people to be happy. → True											
	Characters entertain [audiences] <sub>2</sub> because [they] <sub>1</sub> want people to be happy. → False											
Rel.	The [burst] <sub>1</sub> has been caused by water hammer [pressure] <sub>2</sub> . → Cause-Effect( <i>e</i> <sub>2</sub> , <i>e</i> <sub>1</sub> )											

Table 1: Example sentence, spans, and target label for each task. O = OntoNotes, W = Winograd.

# Contextualized Representation Models

Model	Training objective/task	Structure	Training Corpus	Usage
CoVe	Machine Translation	Two-layer biLSTM	WMT2017: 7 million sentences from web crawl, news, and government proceedings	Top-level activation, concatenated with Glove vectors. (300*2+300 dims)
ELMo	Language Modeling	Two-layer bidirectional LSTM over a context-independent character CNN layer	Billion Word Benchmark	Standard usage
GPT		12-layer Transformer as a left-to-right language model	Toronto Books Corpus	Do not fine-tune, <b>cat(as CoVe)</b> / <b>mix(as ELMo)</b>
BERT	Masked Language Modeling (MLM) & Next Sentence Prediction (NSP)	Deep Transformer (12/24-layer)	Toronto Books Corpus & English Wikipedia	

# Baselines

- **Lexical Baselines**

- Contextualized vs Non-contextualized
- CoVe vs 300 dims GloVe
- ELMo vs context-independent character-CNN of ELMo
- GPT&BERT vs subword embeddings

- **Randomized ELMo**

- Random orthonormal matrices: Invalidate the information of ELMo.

- **Word-Level CNN**

- CNN: See tokens around the center word.



# Results and Analysis

# Results and Analysis


- CoVe, ELMo, GPT, BERT
  - the large gains: mostly syntactic
  - the small gains: mostly semantic
- SRL (core roles) > SRL (non-core roles)
- mix > cat
- Large improvement of BERT on Winograd-style coreference
- GPT trained on BWB  $\approx$  GPT trained on Books Corpus

	CoVe			ELMo			GPT		
	Lex.	Full	Abs. $\Delta$	Lex.	Full	Abs. $\Delta$	Lex.	cat	mix
Part-of-Speech	85.7	94.0	8.4	90.4	<b>96.7</b>	6.3	88.2	94.9	95.0
Constituents	56.1	81.6	25.4	69.1	<b>84.6</b>	15.4	65.1	81.3	<b>84.6</b>
Dependencies	75.0	83.6	8.6	80.4	<b>93.9</b>	13.6	77.7	92.1	<b>94.1</b>
Entities	88.4	90.3	1.9	92.0	<b>95.6</b>	3.5	88.6	92.9	92.5
SRL (all)	59.7	80.4	20.7	74.1	<b>90.1</b>	16.0	67.7	86.0	89.7
Core roles	56.2	81.0	24.7	73.6	<b>92.6</b>	19.0	65.1	88.0	92.0
Non-core roles	67.7	78.8	11.1	75.4	<b>84.1</b>	8.8	73.9	81.3	<b>84.1</b>
OntoNotes coref.	72.9	79.2	6.3	75.3	<u>84.0</u>	8.7	71.8	83.6	<u>86.3</u>
SPR1	73.7	77.1	3.4	80.1	<b>84.8</b>	4.7	79.2	83.5	83.1
SPR2	76.6	80.2	3.6	82.1	83.1	1.0	82.2	<b>83.8</b>	<b>83.5</b>
Winograd coref.	52.1	<b>54.3</b>	2.2	<b>54.3</b>	<u>53.5</u>	-0.8	51.7	52.6	<u>53.8</u>
Rel. (SemEval)	51.0	60.6	9.6	55.7	77.8	22.1	58.2	<b>81.3</b>	<b>81.0</b>
Macro Average	<u>69.1</u>	<u>78.1</u>	9.0	<u>75.4</u>	<u>84.4</u>	9.1	<u>73.0</u>	83.2	<u>84.4</u>

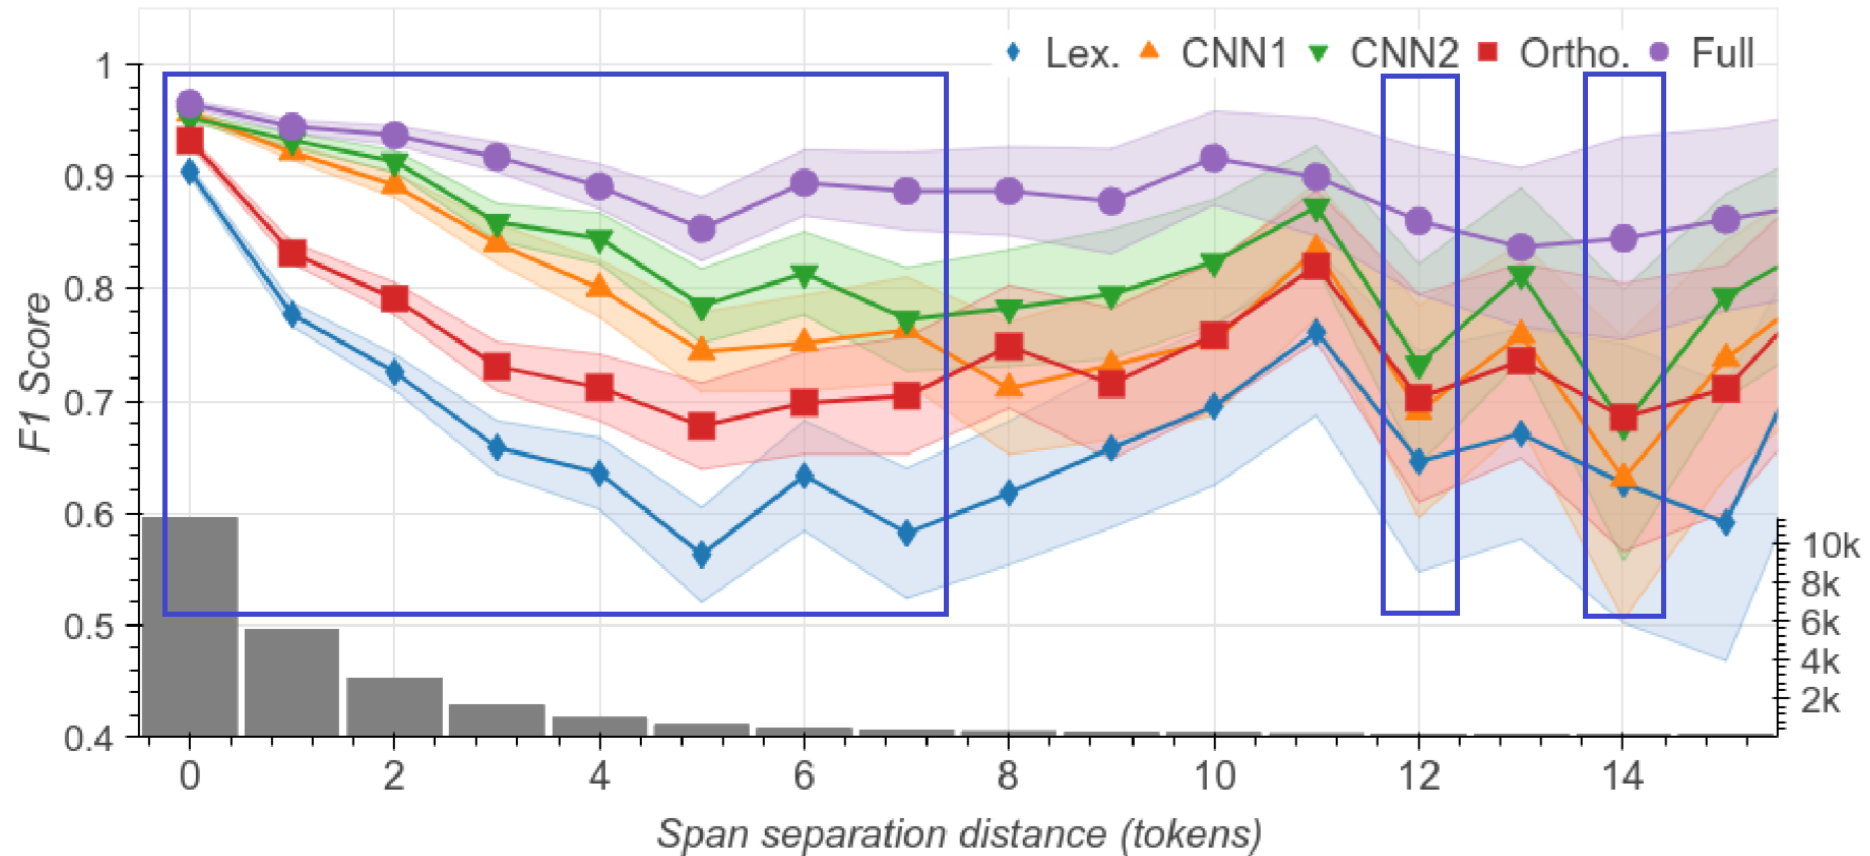
	BERT-base				BERT-large				
	F1 Score			Abs. $\Delta$ ELMo	F1 Score			Abs. $\Delta$ (base)	ELMo
	Lex.	cat	mix		Lex.	cat	mix		
Part-of-Speech	88.4	<b>97.0</b>	96.7	0.0	88.1	96.5	<b>96.9</b>	0.2	0.2
Constituents	68.4	83.7	86.7	2.1	69.0	80.1	<b>87.0</b>	0.4	2.5
Dependencies	80.1	93.0	95.1	1.1	80.2	91.5	<b>95.4</b>	0.3	1.4
Entities	90.9	96.1	96.2	0.6	91.8	96.2	<b>96.5</b>	0.3	0.9
SRL (all)	75.4	89.4	91.3	1.2	76.5	88.2	<b>92.3</b>	1.0	2.2
Core roles	74.9	91.4	93.6	1.0	76.3	89.9	<b>94.6</b>	1.0	2.0
Non-core roles	76.4	84.7	85.9	1.8	76.9	84.1	<b>86.9</b>	1.0	2.8
OntoNotes coref.	74.9	88.7	<u>90.2</u>	6.3	75.7	89.6	<u>91.4</u>	1.2	7.4
SPR1	79.2	84.7	<b>86.1</b>	1.3	79.6	85.1	<b>85.8</b>	-0.3	1.0
SPR2	81.7	83.0	<b>83.8</b>	0.7	81.6	83.2	<b>84.1</b>	0.3	1.0
Winograd coref.	54.3	53.6	<u>54.9</u>	1.4	53.0	53.8	<u>61.4</u>	6.5	7.8
Rel. (SemEval)	57.4	78.3	82.0	4.2	56.2	77.6	<b>82.4</b>	0.5	4.6
Macro Average	75.1	<u>84.8</u>	<u>86.3</u>	1.9	75.2	<u>84.2</u>	<u>87.3</u>	1.08	2.9

# Results and Analysis



- Word-Level CNN closes most of the gap on syntactic tasks.
  - Local information
- The gap between Word-Level and full model is large still on semantic tasks.
  - Long-range information
- Full model > Orthonormal encoder > Lexical baseline
  - A flaw: Why so weird on Winograd coref.?

# Results and Analysis



- Full ELMo holds up better, dropping only 7 F1 score from d=0 to d=8
  - The pre-trained encoder does encode useful long-distance dependencies.

# Results and Analysis – summary

- **Syntactic** or **Semantic**? -> **Syntactic** > **Semantic**
- **Local** or **Long-range**? -> **Local** & **Long-range**
- Different layer, different info.
- Well-designed architecture
- Deep model does well on difficult semantic task

# Discussion

What can we learn from that? &  
What should we do in the future?

# Discussion

- Existing representation models encode more syntactic than semantic info.
  - Bright future: To capture semantic context information
- BERT(-base) outperforms GPT, average 1.6-1.9 F1 score
  - Novel effective training objectives
  - Baidu-ERNIE: 'xxx is one of the top universities in the world.'
- No common sense -> Using knowledgebase
  - Learn semantics and common sense

# References

- Tenney I, Xia P, Chen B, et al. What do you learn from context? Probing for sentence structure in contextualized word representations[J]. 2018.
- McCann B, Bradbury J, Xiong C, et al. Learned in translation: Contextualized word vectors[C]//Advances in Neural Information Processing Systems. 2017: 6294-6305.
- Peters M E, Neumann M, Iyyer M, et al. Deep contextualized word representations[J]. arXiv preprint arXiv:1802.05365, 2018.
- Radford A, Narasimhan K, Salimans T, et al. Improving language understanding by generative pre-training[J]. URL [https://s3-us-west-2.amazonaws.com/openai-assets/research-covers/language-unsupervised/language\\_understanding\\_paper.pdf](https://s3-us-west-2.amazonaws.com/openai-assets/research-covers/language-unsupervised/language_understanding_paper.pdf) , 2018.
- BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding